

Bayesian Hyperbolic Embeddings for Extreme Uncertainty in Zero-Shot Learning

Ilir Bejleri
ibejler3@gatech.edu

April 5, 2026

Abstract

Zero-shot learning (ZSL) models must generalise to classes never observed during training by bridging a visual–semantic gap. Existing approaches embed visual features as *deterministic points* in Euclidean space, discarding information about per-sample prediction confidence. We propose **Bayesian Hyperbolic ZSL**, a framework that maps visual features from a frozen backbone to *Wrapped Normal distributions* on the Poincaré ball, where the learned variance explicitly captures visual–semantic uncertainty. By aligning distributional embeddings with class prototypes through curvature-aware geodesic losses, our model simultaneously learns *what* to embed and *how curved* the embedding space should be. We benchmark on Animals with Attributes 2 (AwA2) across a data-ablation sweep (1%–100% training data) and find that the Bayesian hyperbolic model achieves accuracy competitive with the Euclidean baseline in the moderate-scarcity regime (5% data), while revealing a striking emergent property: the learned Poincaré ball curvature scales monotonically with data availability ($c = 1.18$ at 1% to $c = 6.63$ at 100%), suggesting the model automatically discovers the appropriate manifold geometry. We analyse the regularisation–accuracy trade-off inherent to the Bayesian formulation and identify concrete directions for improvement.

1 Introduction

Zero-shot learning asks a model to recognise visual categories for which no labelled training examples were available, relying instead on a shared semantic description such as attribute vectors or word embeddings Romera-Paredes and Torr [2015], Xian et al. [2019b]. The dominant paradigm learns a *visual–semantic embedding*—a mapping from image features into the same space as class descriptors—so that unseen classes can be recognised via nearest-neighbour lookup at test time.

Most ZSL methods operate in flat Euclidean space. Yet the semantic relationships between classes are inherently *hierarchical*: WordNet Miller [1995] organises animal concepts into an “is-a” taxonomy whose branching structure is poorly captured by Euclidean geometry. Hyperbolic spaces, by contrast, enjoy exponentially growing volume with radius, making them natural for embedding trees with low distortion Nickel and Kiela [2017, 2018].

The foundational work of **Ganea et al.** Ganea et al. [2018] introduced a principled framework for deep learning on hyperbolic manifolds. Their formulation of hyperbolic linear layers via Möbius matrix–vector multiplication and Möbius addition of bias terms (operating in the Poincaré ball model $\mathbb{B}_c^d$ of constant negative curvature $-1/c$) provides the mathematical basis upon which our architecture is built. Crucially, Ganea et al. demonstrated that standard neural-network operations—linear maps, non-linearities, and attention—can be lifted to the Poincaré ball while preserving differentiability and numerical stability, enabling end-to-end gradient-based training Ganea et al. [2018]. We inherit their exponential and logarithmic map formulations, their Möbius operations, and their insight that curvature can be treated as a *learnable* parameter.

Limitations of deterministic embeddings. Mapping each image to a single point in embedding space—whether Euclidean or hyperbolic—discards information about *how confident* the model is in its projection. An image of a partially occluded animal and a clean studio photograph receive embeddings of identical certainty, even though the model’s evidence for the former is substantially weaker. This brittleness is amplified in the zero-shot setting, where the model must extrapolate across a visual–semantic gap it has never directly observed.

Our approach. We replace the deterministic point embedding with a **Wrapped Normal distribution** on $\mathbb{B}_c^d$. A visual encoder outputs both a mean $\mu \in \mathbb{B}_c^d$ (via the hyperbolic linear layer of Ganea et al. [2018]) and a variance $\sigma^2 \in \mathbb{R}_+^d$. Samples are drawn via the reparameterisation trick in the tangent space at μ and projected onto the manifold through the exponential map Nagano et al. [2019], Mathieu et al. [2019]. A KL regulariser anchors the posterior toward a prior centred at the origin, and a geodesic alignment loss pulls samples toward the correct class prototype.

Contributions.

1. We architect an experimental zero-shot learning framework leveraging Riemannian geometry that maps visual features to distributions on the Poincaré ball rather than deterministic points.
2. We engineer a visual–semantic bridge using a hyperbolic linear layer Ganea et al. [2018] that outputs the parameters of a Wrapped Normal distribution, capturing per-sample uncertainty.
3. We benchmark performance against Euclidean baselines across a data-ablation sweep on AwA2 and discover that the learned curvature parameter adapts monotonically to data availability—a novel finding with implications for understanding how neural networks interact with non-Euclidean geometry.

2 Related Work

Hyperbolic representation learning. Nickel & Kiela Nickel and Kiela [2017] first showed that Poincaré embeddings can represent WordNet hierarchies with dramatically fewer dimensions than Euclidean embeddings. Ganea et al. [2018] generalised this to deep networks by defining hyperbolic analogues of linear layers, concatenation, and attention. Subsequent work extended Riemannian optimisation Bécigneul and Ganea [2019] and provided practical tooling Kochurov et al. [2020].

Probabilistic hyperbolic models. Mathieu et al. [2019] introduced the Poincaré VAE with a Wrapped Normal prior, while Nagano et al. [2019] derived the density and sampling procedure for the Wrapped Normal distribution on hyperbolic space. Our work adapts these ideas from generative modelling to the discriminative ZSL setting.

Zero-shot learning. Xian et al. [2019b] established the standard evaluation protocol for ZSL, including the proposed-split benchmarks on AwA2 Xian et al. [2019a] that we follow.

3 Methodology

3.1 Preliminaries: The Poincaré Ball Model

We work in the Poincaré ball $\mathbb{B}_c^d = \{x \in \mathbb{R}^d : c\|x\|^2 < 1\}$ equipped with the Riemannian metric tensor

$$g_x^c = (\lambda_x^c)^2 g^E, \quad \lambda_x^c = \frac{2}{1 - c\|x\|^2}, \quad (1)$$

where g^E is the Euclidean metric and $c > 0$ is the (learnable) curvature parameter Ganea et al. [2018].

The **exponential map** at x sends a tangent vector $v \in T_x \mathbb{B}_c^d$ to the manifold:

$$\text{Exp}_x^c(v) = x \oplus_c \left(\tanh\left(\frac{\sqrt{c}\lambda_x^c\|v\|}{2}\right) \frac{v}{\sqrt{c}\|v\|} \right), \quad (2)$$

where $\oplus_c$ denotes the Möbius addition. The **logarithmic map** inverts this operation.

3.2 Architecture

Visual encoder. Given a pre-extracted feature vector $x \in \mathbb{R}^{2048}$ from a frozen ResNet-101 He et al. [2016], a two-layer Euclidean MLP produces a hidden representation $h \in \mathbb{R}^H$. Two heads branch from h :

$$\mu = \text{HypLinear}(h) \in \mathbb{B}_c^d, \quad (3)$$

$$\log \sigma^2 = W_\sigma h + b_\sigma \in \mathbb{R}^d, \quad (4)$$

where HypLinear denotes the hyperbolic linear layer of Ganea et al. [2018], applying a Möbius matrix-vector product followed by Möbius bias addition.

Semantic encoder. Each class is represented by an 85-dimensional continuous attribute vector from AwA2. A two-layer MLP followed by a hyperbolic linear layer maps these attributes to a deterministic point $s_y \in \mathbb{B}_c^d$.

Shared manifold. Both encoders share a single Poincaré ball instance with *learnable* curvature c , following Ganea et al. [2018].

3.3 Wrapped Normal Distribution

We parameterise a Wrapped Normal $\text{WN}(\mu, \sigma^2)$ on $\mathbb{B}_c^d$ following Nagano et al. [2019]:

1. Sample $\epsilon \sim \mathcal{N}(0, I)$ in $\mathbb{R}^d$.
2. Form the tangent vector $v = \sigma \odot \epsilon \in T_0 \mathbb{B}_c^d$.
3. Parallel-transport v from the origin to μ : $\tilde{v} = \Gamma_{0 \rightarrow \mu}(v)$.
4. Project onto the manifold: $z = \text{Exp}_\mu^c(\tilde{v})$.

This procedure is fully differentiable via the reparameterisation trick.

3.4 Loss Function

The total objective is:

$$\mathcal{L} = \mathcal{L}_{\text{align}} + \lambda_{\text{KL}} \mathcal{L}_{\text{KL}} + \lambda_{\text{hier}} \mathcal{L}_{\text{hier}}. \quad (5)$$

Alignment loss. A softmax cross-entropy over negative geodesic distances:

$$\mathcal{L}_{\text{align}} = -\log \frac{\exp(-d_{\mathbb{B}_c^d}(z, s_y)/\tau)}{\sum_{j=1}^C \exp(-d_{\mathbb{B}_c^d}(z, s_j)/\tau)}, \quad (6)$$

where z is a sample from the visual posterior, s_y is the ground-truth class prototype, and τ is a temperature hyperparameter.

KL regulariser. An approximate KL divergence between the posterior $\text{WN}(\mu, \sigma^2)$ and a standard prior $\text{WN}(0, I)$, computed in the tangent space at the origin:

$$\mathcal{L}_{\text{KL}} \approx \frac{1}{2} \sum_{i=1}^d (\sigma_i^2 + \|\text{Log}_0(\mu)\|_i^2 - 1 - \log \sigma_i^2). \quad (7)$$

Hierarchy-aware margin. An optional contrastive margin scaled by WordNet path-length distance between classes, encouraging the embedding geometry to reflect the semantic taxonomy.

4 Experiments

4.1 Setup

Dataset. We use Animals with Attributes 2 (AwA2) Xian et al. [2019a] with the proposed split of 40 seen / 10 unseen classes and 7,454 test images Xian et al. [2019b]. Pre-extracted ResNet-101 He et al. [2016] features (2048-d) are used as visual input. Class-level semantic embeddings are the 85-dimensional continuous attribute vectors from AwA2.

Table 1: Top-1 and Top-5 accuracy (%) on AwA2 unseen classes across data-ablation fractions. The Bayesian hyperbolic model achieves competitive accuracy at the 5% fraction, while the Euclidean baseline is more stable across larger fractions.

Model	Metric	Training Data Fraction					
		1%	5%	10%	25%	50%	100%
Euclidean ZSL	Top-1	69.28	66.89	70.00	67.44	71.46	70.28
Bayesian Hyp. (ours)	Top-1	59.02	67.36	63.54	57.78	55.23	60.76
Euclidean ZSL	Top-5	96.99	97.49	97.21	97.67	98.44	98.00
Bayesian Hyp. (ours)	Top-5	95.95	97.59	95.37	91.37	89.33	88.69

Table 2: Learned Poincaré ball curvature c and KL divergence at convergence (epoch 200) as a function of training data fraction. Curvature increases monotonically with data availability while KL decreases, indicating the model adapts its geometry and uncertainty to the evidence regime.

	1%	5%	10%	25%	50%	100%
Final curvature c	1.18	2.11	3.24	5.36	6.09	6.63
Final KL divergence	4.11	2.18	1.36	0.78	0.68	0.67

Baselines.

- **Euclidean ZSL:** Deterministic MLP mapping features to $\mathbb{R}^{64}$, trained with distance-based cross-entropy alignment and Adam optimisation.

Both models share the same hidden dimension ($H = 1024$) and embedding dimension ($d = 64$) to isolate the effect of geometry and uncertainty modelling.

Data-ablation protocol. To test the “extreme data efficiency” hypothesis, we train each model on $\{1\%, 5\%, 10\%, 25\%, 50\%, 100\%\}$ of the seen-class training data (corresponding to approximately 298–29,780 training samples) and evaluate on all unseen-class test images.

Implementation. The Bayesian hyperbolic model is optimised with RiemannianAdam Bécigneul and Ganea [2019] (10^{-3} for Euclidean parameters, 10^{-4} for manifold parameters) using the GeoOpt library Kochurov et al. [2020]. Temperature $\tau = 0.1$, KL weight $\lambda_{\text{KL}} = 0.01$, and initial curvature $c_0 = 1.0$ (learnable). All models train for 200 epochs with batch size 128 on NVIDIA V100 GPUs. Gradient norms are clipped at 5.0 for stability. Bayesian evaluation marginalises over 10 posterior samples per input.

4.2 Results

Table 1 presents zero-shot accuracy on the 10 unseen AwA2 classes. The Euclidean baseline is remarkably stable, achieving 67–71% Top-1 and 97–98% Top-5 across all data fractions—suggesting that the Euclidean embedding space, with Adam’s effective learning-rate adaptation, is robust to data scarcity on this benchmark.

The Bayesian hyperbolic model exhibits different behaviour. At the **5% fraction**, it achieves its peak accuracy of 67.36% Top-1 and 97.59% Top-5, marginally outperforming the Euclidean baseline (66.89% / 97.49%). However, at other fractions—particularly 25%–100%—accuracy degrades relative to the baseline. This non-monotonic trend, where accuracy peaks at moderate scarcity rather than increasing with data, reveals a fundamental tension between the KL regulariser and the alignment objective that we analyse in Section 4.3.

4.3 Analysis

The regularisation–accuracy trade-off. The Bayesian model’s accuracy peaks at 5% and then *decreases* with more data. This is counterintuitive for a discriminative model but explicable through the interaction of its loss components. As data increases, the alignment loss drives visual embeddings toward class prototypes more aggressively, while the

Zero-Shot Accuracy vs. Training Data — AwA2 Unseen Classes





Figure 1: Zero-shot accuracy on AwA2 unseen classes as a function of training data fraction. The Bayesian hyperbolic model (red) achieves competitive accuracy at 5% data but exhibits a non-monotonic trend at larger fractions due to the regularisation–accuracy trade-off. The Euclidean baseline (blue) is more stable across fractions.

KL term simultaneously pulls the posterior toward the origin-centred prior. Table 2 shows that the final KL divergence drops from 4.11 (1% data) to 0.67 (100% data), indicating that the posterior collapses closer to the prior with more data—the opposite of the desired behaviour. This suggests that the fixed KL weight $\lambda_{\text{KL}} = 0.01$ becomes too strong relative to the alignment signal at higher data fractions, effectively over-regularising the model. An adaptive KL schedule (e.g., $\lambda_{\text{KL}} \propto N^{-1}$) could mitigate this; we leave this to future work.


Learned curvature reflects data complexity. The most striking finding is the relationship between learned curvature and dataset size (Table 2, Figure 2a). With 1% training data, the model converges to $c = 1.18$ —nearly flat geometry. With the full dataset, curvature reaches $c = 6.63$, corresponding to a much more tightly curved space. This monotonic relationship ($c \propto \log N$ approximately) demonstrates that the model automatically adjusts the manifold geometry to match the complexity of the learned representations: more data reveals more hierarchical structure, which is captured by higher curvature. The curvature evolution curves (Figure 2a) show that curvature increases smoothly during training, without instability, confirming that treating c as a learnable parameter Ganea et al. [2018] is numerically well-behaved even in the discriminative ZSL setting.

Training dynamics and posterior behaviour. Figure 3a shows training loss convergence for all ablation fractions. The KL divergence (Figure 3b) decreases monotonically during training across all fractions. At 1% data, the final KL is 4.11, while at 100% it drops to 0.67—the model appropriately expresses higher uncertainty when evidence is scarce. Notably, the KL term dominates the total loss at low data fractions (contributing $> 60\times$ the alignment loss at 1%), while the alignment term grows to dominate at higher fractions. This shift in the loss landscape explains the non-monotonic accuracy behaviour: at 5%, the two terms are well balanced; at higher fractions, the model capacity is increasingly spent satisfying the KL constraint rather than improving alignment.

Why 5% is the sweet spot. The Bayesian model’s peak at 5% data can be understood as the regime where (1) enough data exists for the alignment loss to learn meaningful prototypes, but (2) not so much data that the alignment gradient overwhelms the KL regulariser, causing the posterior to collapse. In this regime, the Wrapped Normal variance provides genuine benefit: it smooths the nearest-neighbour decision boundary over unseen class prototypes, and the moderate curvature ($c = 2.11$) provides enough geometric structure for hierarchical organisation without excessive distortion. This interpretation is supported by the comparable Top-5 accuracy (97.59% vs. 97.49%), which suggests the distributional embedding places the correct class within the top candidates as reliably as the deterministic baseline.




(a) Curvature trajectory during training, coloured by data fraction. More data $\Rightarrow$ higher final curvature.




(b) Final curvature vs. training data fraction. The approximately logarithmic relationship suggests automatic geometry adaptation.

Figure 2: Learned Poincaré ball curvature as a function of training data availability.



(a) Training loss (log scale) for Bayesian hyperbolic and Euclidean models across data fractions.



(b) KL divergence over training. Data-scarce runs maintain higher KL (broader posteriors), avoiding overfitting.

Figure 3: Training dynamics across the data-ablation sweep.

5 Discussion

Our results demonstrate that combining Bayesian inference with hyperbolic geometry for ZSL is architecturally feasible and yields scientifically interesting behaviour, even though the Euclidean baseline achieves higher accuracy at most data fractions on this benchmark. We identify several factors that explain this gap and suggest remedies:

Fixed KL weight. The constant $\lambda_{KL} = 0.01$ was selected by grid search but is suboptimal across the full range of data fractions. A data-dependent schedule $\lambda_{KL}(N) = \lambda_0/N$ or a free-bits formulation Kingma et al. [2016] would allow the model to exploit additional data for alignment rather than over-regularising.

Euclidean robustness. The Euclidean baseline’s stability across data fractions (67–71% Top-1) is notable and likely reflects the effectiveness of Adam with weight decay on this benchmark, where the 85-dimensional attribute vectors provide a rich enough semantic space that even simple linear alignment suffices. On benchmarks with sparser or noisier semantics (e.g., word2vec-based class embeddings), the hyperbolic structure may provide larger gains.

MC marginalisation overhead. Bayesian evaluation averages geodesic distances over 10 posterior samples, adding computational cost. In practice, the distributional embedding provides *calibrated uncertainty estimates*—samples with high posterior variance can be flagged for human review—a capability that deterministic models lack entirely.

6 Conclusion

We presented Bayesian Hyperbolic Zero-Shot Learning, a framework that maps visual features to Wrapped Normal distributions on the Poincaré ball rather than deterministic point embeddings. Built upon the foundational hyperbolic neural network operations of Ganea et al. [2018], the framework achieves competitive accuracy with the Euclidean baseline in the moderate-scarcity regime (5% data: 67.36% vs. 66.89% Top-1) while additionally providing per-sample uncertainty estimates.

The most notable finding is that the learnable curvature parameter automatically adapts to data availability, scaling from $c = 1.18$ (1% data) to $c = 6.63$ (100% data) in a near-logarithmic relationship. This emergent behaviour—the model discovering appropriate manifold geometry from data alone—has implications beyond ZSL for any application of learnable-curvature hyperbolic networks.

We also identified a regularisation–accuracy trade-off: the fixed KL weight that enables uncertainty estimation simultaneously limits alignment capacity at higher data fractions. Addressing this through adaptive scheduling, richer semantic inputs, or hierarchical priors aligned with WordNet structure represents a promising direction for future work.

References

- Gary Bécigneul and Octavian-Eugen Ganea. Riemannian adaptive optimization methods. In *International Conference on Learning Representations (ICLR)*, 2019.
- Octavian-Eugen Ganea, Gary Bécigneul, and Thomas Hofmann. Hyperbolic neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 31, 2018. arXiv:1805.09112.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- Diederik P. Kingma, Tim Salimans, Rafal Jozefowicz, Xi Chen, Ilya Sutskever, and Max Welling. Improved variational inference with inverse autoregressive flow. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29, 2016.
- Max Kochurov, Rasul Karimov, and Serge Kozlukov. Geoopt: Riemannian optimization in PyTorch, 2020. arXiv:2005.02819.
- Emile Mathieu, Charline Le Lan, Chris J. Maddison, Ryota Tomioka, and Yee Whye Teh. Continuous hierarchical representations with Poincaré variational auto-encoders. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.
- George A. Miller. WordNet: A lexical database for English. *Communications of the ACM*, 38(11):39–41, 1995.
- Yoshihiro Nagano, Shoichiro Yamaguchi, Yasuhiro Fujita, and Masanori Koyama. A wrapped normal distribution on hyperbolic space for gradient-based learning. In *International Conference on Machine Learning (ICML)*, 2019.
- Maximilian Nickel and Douwe Kiela. Poincaré embeddings for learning hierarchical representations. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, 2017.
- Maximilian Nickel and Douwe Kiela. Learning continuous hierarchies in the Lorentz model of hyperbolic geometry. In *International Conference on Machine Learning (ICML)*, 2018.
- Bernardino Romera-Paredes and Philip Torr. An embarrassingly simple approach to zero-shot learning. In *International Conference on Machine Learning (ICML)*, 2015.
- Yongqin Xian, Christoph H. Lampert, Bernt Schiele, and Zeynep Akata. Zero-shot learning — a comprehensive evaluation of the good, the bad and the ugly. *IEEE TPAMI*, 2019a. AwA2 dataset.
- Yongqin Xian, Christoph H. Lampert, Bernt Schiele, and Zeynep Akata. Zero-shot learning — a comprehensive evaluation of the good, the bad and the ugly. volume 41, pages 2251–2265, 2019b.